

Collision Avoidance using Deep Q-Network


Robot Learning Project

Chethan Mysore Parameshwara

Outline

- > When to use DQN ? - Problem Formulation
- > Why DQN? Why not Q-Learning ?
- > What is DQN ?
- > How to train DQN ?
- > Does it really work ? Demo
- > Who helped me ? References

When to use DQN ? - Problem Formulation



> Markov Random Process

probability of the next state s_{i+1} depends only on current state s_i and action a_i , but not on preceding states or actions.

> Episode - $s_0, a_0, r_1, s_1, a_1, r_2, s_2, \dots, s_{n-1}, a_{n-1}, r_n, s_n$

> Policy - $\pi = P(a/s)$

> Rewards -

```
if not Game_Over:
    if action == FORWARD:
        reward = 5
    else:
        reward = 1
else:
    reward = -200
```

Why DQN ? Why not Q-Learning ?

We know that

$$R_t = r_t + \gamma(r_{t+1} + \gamma(r_{t+2} + \dots)) = r_t + \gamma R_{t+1}$$

$$Q(s_t, a_t) = \max R_{t+1}$$

$$\pi(s) = \operatorname{argmax}_a Q(s, a)$$

We iteratively approximate the Q-function using the Bellman equation

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

Q - Learning

$$Q[s, a] = Q[s, a] + \alpha (r + \gamma \max Q[s, a] - Q[s', a'])$$

Q - Learning Limitations

- > Multi-dimensional state
- > Multiple action
- > Complexity in searching Q - table

What is DQN ?

- > Approximate Q-function with a neural network
- > Extracts features from multidimensional data
- > Loss Function

$$L = \frac{1}{2} \left[\underbrace{r + \max_{a'} Q(s', a')}_{\text{target}} - \underbrace{Q(s, a)}_{\text{prediction}} \right]^2$$

Problems with DQN

- > correlations present in the sequence of observations
- > correlations between the action-values (Q) and the target values ($r + \max (Q)$)

Workaround

- > Experience Replay
- > Fixed Target Q-Network

How to train DQN ?

Network Architecture

- Input Size - 100
- Output Size - 3
- Network = 3 Fully Connected Layer
- Activation Function - ReLU
- Optimization Solver- RMSprop

Network Parameters

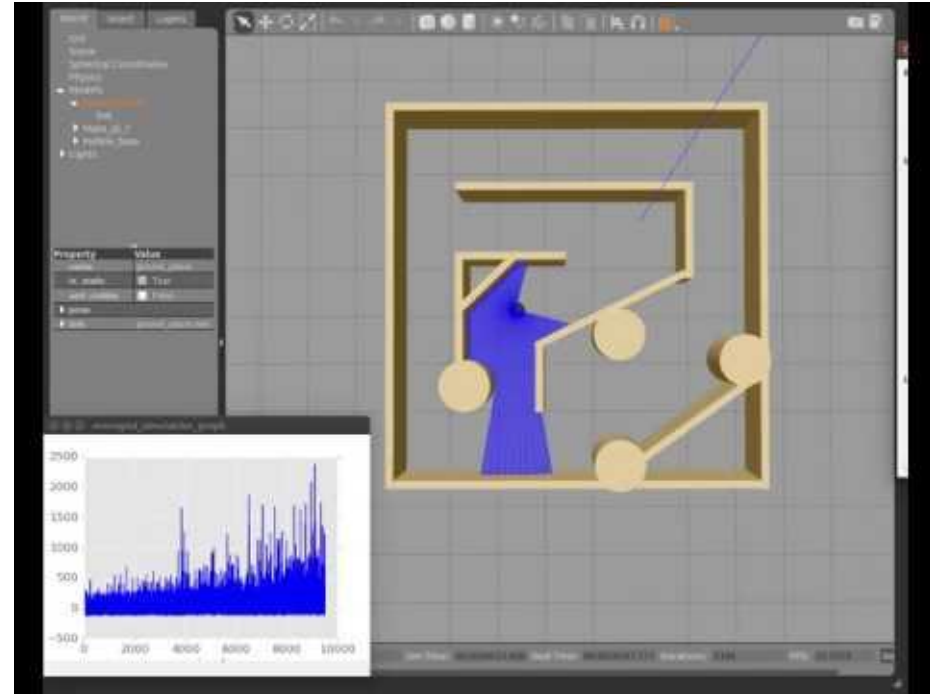
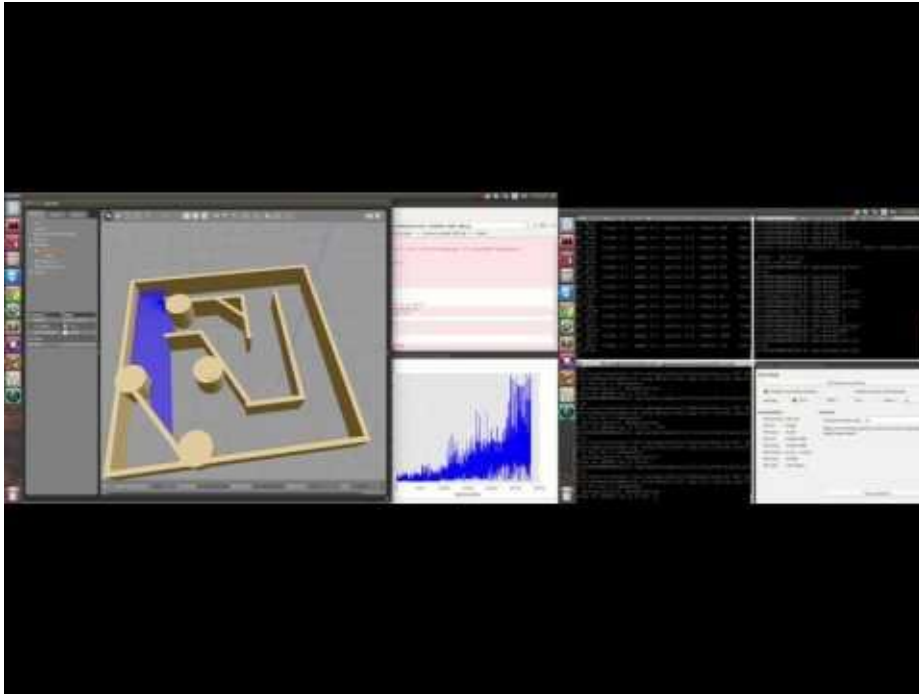
- Max_Episodes = 15000
- Max_steps_episode = 1000
- Exploration_Rate = 1 (adaptive)
- Minibatch_size = 64
- Learning_Rate = 0.8
- Discount_Factor = 0.99

Hardware - NVIDIA GF106GL [Quadro 2000] GPU (Outdated !!)

Software - Keras, ROS, Gazebo

Programming Language - Python

Does it really work ? Demo



Source Code - <https://github.com/analogicalnexus/gym-gazebo>

Who helped me ? References

- Zamora, Iker, et al. "Extending the OpenAI Gym for robotics: a toolkit for reinforcement learning using ROS and Gazebo." *arXiv preprint arXiv:1608.05742* (2016).
- Mnih, Volodymyr, et al. "Human-level control through deep reinforcement learning." *Nature* 518.7540 (2015): 529-533.
- <https://www.nervanasys.com/demystifying-deep-reinforcement-learning/>

Thank you